

3/30

A quick test drive of randomForest on a smallish set (200 training) does OK but seems to have some degree of biased error.

High var explained (97%).

compared to a (v. rudimentary) linear model, there's less (obvious) structured error.

→ IDK what would happen predicting far out of sample.

Looking at some 1:1 plots



some are great and some are systematically not.

→ Is that something that would be improved with a) more values or b) a more nuanced algorithm? (e.g. NN instead of RF?)